TD1/5: Project

The goal of this project is to predict *part_assiette_chomage_partiel* one month ahead. You will be guided through the differents steps that are fundamental in a Data Science project. You still have an important part of autonomy in your code, methods, analysis and decisions. Several interpretations often coexist as long as they are coherent. Do your best, be courageous and deploy your Data Science artillery!

Exercise 1: Set up

1. Load data:
   masse-salariale-et-assiette-chomage-partiel-mensuelles-du-secteur-prive_modif.csv[1]
2. See number of samples (rows) and features (columns)
3. See data type
4. Set *dernier_jour_du_mois* as index
5. Cast index as datetime
6. Sort index in ascending order

Exercise 2: Data Analysis

1. Produce standard descriptive statistics
2. Visualize data (many plots can be done in one line of code using Pandas and Seaborn)

Exercise 3: Data Cleaning

1. Check for missing values (some are more subtle than a explicit NaN)
2. Impute these missing values with at least 2 methods seen in the lectures, don't delete them in this project (we need all the dates in a Times Series problem)
3. Check and treat outlier(s)

Exercise 4: Feature Engineering

1. Add a feature *is_year_end*
   - 1 when month is november or december
   - 0 otherwise

[1] Data is a modified version from this source

Exercise 5: Prediction

1. Split your data into a train set (60% of data) and a test set (40%)
2. Use a linear regression to predcit *part_assiette_chomage_partiel* 1 month ahead
    - you should shift your features (in time) compared to your target
    - find tutorials, there are a lot of them, its the only way toward autonomous learning!
3. How good is your prediction?
    - Plot the predicted values on the same graph than the actual values
    - Use metric(s) to evaluate your model on both the train and test sets
    - Interpret the results
    - Give advices to your (hypothetical) colleague to continue your work

Exercise 5:.1 Bonus

1. Make a prediction without the added variable *is_year_end*, what is the impact?
2. Use a Ridge regression in place of the Linear regression (you might become happy about the results!)
3. Use a **polynomial** regression to predcit 1 month ahead (find tutorials, there are a lot of them, and its the only way to learn autonomously!)
4. Predict 2 months ahead, then 3 and 4 months ahead. If your code is written correctly, it should only require to manually change the value of a constant.