

Introduction to Data Science

with Python

Alexis Bogroff

May 24, 2022



Alexis Bogroff

Lecturer and Mentor in Data Science
at Paris 1 Panthéon-Sorbonne, ESILV,
Openclassrooms, EM-Lyon



Alexis Bogroff

Lecturer and Mentor in Data Science
at Paris 1 Panthéon-Sorbonne, ESILV,
Openclassrooms, EM-Lyon

- 4 years Teaching Assistant and lecturer in VBA, Python for finance, SQL, Data Analysis and Data Science
- 9 months Researcher Assistant at Paris 1 Panthéon-Sorbonne within H2020 European Project
- 1 year Data Scientist at Pléiade Asset Management

- Centrality:
 - Goal: representation of the majority's value
 - Mean (average): average age, mean size
 - Median: median salary, median patrimony
- Dispersion:
 - Goal: majority's spread (variation) around the central value
 - Standard deviation (sqrt variance): financial markets volatility
 - Interquartile Range (IQR)
 - Min-Max: job proposal salary

Pandas - Mean

- Sensible to extreme values
 - Age: good representation
 - Patrimony: biased, not representative

$$\text{mean} = \frac{x_1 + x_2 + \dots + x_n}{n}$$

	name	age	group	country	patrimony
0	ju	25	1	fr	3000
1	ma	36	1	gb	7000
2	lo	40	2	gb	2000
3	fi	18	3	fr	60000
4	na	25	3	es	8000000
5	pe	40	3	it	4000

```
mean_age = df['age'].mean()
int(mean_age)

✓ 0.2s

30

mean_patrimony = df['patrimony'].mean()
int(mean_patrimony)

✓ 0.3s

1346000
```

Pandas - Median

- Insensitive to extreme values
 - Age: good representation
 - Patrimony: good representation of the majority

Order values, then:

$$\text{median} = \frac{X_{center_2} - X_{center_1}}{2}$$

	name	age	group	country	patrimony
0	ju	25	1	fr	3000
1	ma	36	1	gb	7000
2	lo	40	2	gb	2000
3	fi	18	3	fr	60000
4	na	25	3	es	8000000
5	pe	40	3	it	4000

```
median_age = df['age'].median()
int(median_age)

✓ 0.2s

30

median_patrimony = df['patrimony'].median()
int(median_patrimony)

✓ 0.2s

5500
```

Pandas - Standard Deviation (std)

- Sensible to extreme values
- In the unit of the variable
- Interpretable
 - Age: good representation
 - Patrimony: Patrimony: biased, not representative

$$\bar{x} = \text{mean}(\text{values})$$

$$\text{std} = \sqrt{\frac{(x_1 - \bar{x})^2 + \dots + (x_n - \bar{x})^2}{n}}$$

	name	age	group	country	patrimony
0	ju	25	1	fr	3000
1	ma	36	1	gb	7000
2	lo	40	2	gb	2000
3	fi	18	3	fr	60000
4	na	25	3	es	8000000
5	pe	40	3	it	4000

```
median_age = df['age'].median()  
int(median_age)
```

✓ 0.2s

30

```
median_patrimony = df['patrimony'].median()  
int(median_patrimony)
```

✓ 0.2s

5500

Pandas - Interquartile Range (iqr)

- Insensitive to extreme values
- In the unit of the variable
- Interpretable
 - Age: good representation
 - Patrimony: quite good representation

$$iqr = 3rd_quantile - 1st_quantile$$

	name	age	patrimony
0	ma	36	7000
1	lo	40	2000
2	fi	18	5000
3	na	25	8000000
4	pe	40	4000

```
iqr_age = df_temp['age'].quantile(.75) \
|         - df_temp['age'].quantile(.25)
int(iqr_age)
✓ 0.2s

15

iqr_patrimony = df_temp['patrimony'].quantile(.75) \
|             - df_temp['patrimony'].quantile(.25)
int(iqr_patrimony)
✓ 0.5s

1

3000

# What is a quantile, example:
# These are exact values when 25% and 75% points exist
# Otherwise, linear interpolation
print(int(df_temp['patrimony'].quantile(.25)))
print(int(df_temp['patrimony'].quantile(.75)))
✓ 0.4s

4000
7000
```

¹Quantile values 1st: 25%, 2nd: 50%, 3rd: 75% - after ordering

Pandas - Min Max

- Sensible to extreme values
- In the unit of the variable
- Interpretable
- Easy to compute
- Idea of max range

$$\text{min-max} = \text{max}(\text{values}) - \text{min}(\text{values})$$

	name	age	patrimony
0	ma	36	7000
1	lo	40	2000
2	fi	18	5000
3	na	25	8000000
4	pe	40	4000

```
min_max_age = df['age'].max() - df['age'].min()
min_max_age
✓ 0.2s
22


min_max_patrimony = df['patrimony'].max() \
- df['patrimony'].min()
min_max_patrimony
✓ 0.2s
7998000
```

- Univariate Analysis
 - Time Series
 - Trend
 - Seasonality
 - Auto-correlation
 - Other quantitative variables
 - Qualitative variables
- Multivariate Analysis (between variables)
 - Quantitative variables
 - Linear
 - Non-Linear
 - Qualitative variables

Univariate Analysis

- Time Series
 - Generate plot using Pandas DataFrame method


```
df_time.plot()  
✓ 0.1s  
<AxesSubplot:xlabel='date'>
```



revenue	
date	
2021-01-01	12000
2021-02-01	13000
2021-03-01	13000
2021-04-01	13000
2021-05-01	18000
2021-06-01	19000
2021-07-01	21000
2021-08-01	23000
2021-09-01	18000
2021-10-01	13000
2021-11-01	13000
2021-12-01	13000
2022-01-01	12000
2022-02-01	13000
2022-03-01	13000
2022-04-01	13000
2022-05-01	18000
2022-06-01	19000
2022-07-01	21000
2022-08-01	23000
2022-09-01	20000
2022-10-01	20000
2022-11-01	21000
2022-12-01	20000

Univariate Analysis

- Time Series
 - Compute overall trend




```
revenue_first = df_time['revenue'].iloc[0]
revenue_last = df_time['revenue'].iloc[-1]
n_dates = len(df_time)
(revenue_last - revenue_first) / n_dates
```

✓ 0.2s

333.3333333333333

Univariate Analysis


- Time Series
 - Compute rolling mean
 - Get trend along time
 - Smoothen curve, easier to read
 - + info vs overall trend
 - Delay first dates



```
N_MONTHS = 6  
rolling_mean = df_time['revenue'].rolling(window=N_MONTHS).mean()  
rolling_mean.plot()
```


✓ 0.1s

<AxesSubplot:xlabel='date'>



Univariate Analysis

- Time Series
 - Autocorrelogram
 - Correlation between dates (lags)




```
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

✓ 0.4s

```
plot_acf(df_time['revenue'])
```

None

✓ 0.6s



Univariate Analysis

- Other quantitative variables
 - Descriptive statistics

	transaction	price
0	transac_olp	15
1	transac_ixh	17
2	transac_qkh	20
3	transac_qlz	20
4	transac_mal	19
5	transac_fjh	17
6	transac_rdn	20
7	transac_oaj	20
8	transac_taz	23
9	transac_tgs	17

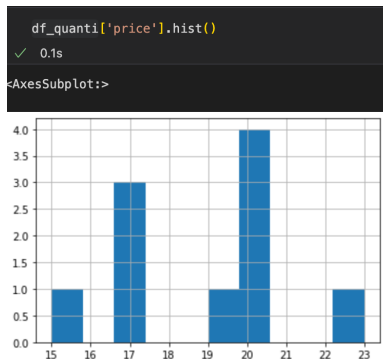
```
df_quant1['price'].describe()
✓ 0.2s
```

count	10.000000
mean	18.800000
std	2.299758
min	15.000000
25%	17.000000
50%	19.500000
75%	20.000000
max	23.000000
Name: price, dtype: float64	

Univariate Analysis

- Other quantitative variables
 - Histogram
 - Overview
 - Extreme values
 - All information

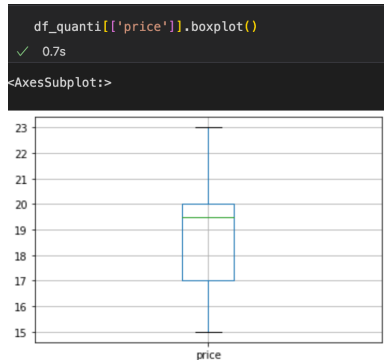
	transaction	price
0	transac_olp	15
1	transac_ixh	17
2	transac_qkh	20
3	transac_qlz	20
4	transac_mal	19
5	transac_fjh	17
6	transac_rdn	20
7	transac_oaj	20
8	transac_taz	23
9	transac_tgs	17



Univariate Analysis

- Other quantitative variables
 - Boxplot
 - Overview
 - Extreme values
 - Contracted information
 - Quartiles
 - Outliers (seaborn)

	transaction	price
0	transac_olp	15
1	transac_ixh	17
2	transac_qkh	20
3	transac_qlz	20
4	transac_mal	19
5	transac_fjh	17
6	transac_rdn	20
7	transac_oaj	20
8	transac_taz	23
9	transac_tgs	17



Univariate Analysis

- Qualitative variables
 - Information on categories

	transaction	client_type
0	transac_olp	veggie
1	transac_ixh	veggie
2	transac_qkh	vegan
3	transac_qlz	omnivorous
4	transac_mal	vegan
5	transac_fjh	veggie
6	transac_rdn	veggie
7	transac_oaj	omnivorous
8	transac_taz	vegan
9	transac_tgs	veggie

```
print(df_quali['client_type'].nunique())  
print(df_quali['client_type'].unique())
```

✓ 0.3s

3

['veggie' 'vegan' 'omnivorous']

```
counts = df_quali['client_type'].value_counts()  
counts
```

✓ 0.3s

veggie 5

vegan 3


omnivorous 2

Name: client_type, dtype: int64

Univariate Analysis

- Qualitative variables
 - Bar plot
 - Overview
 - Ordinal variables:
(small, medium,
large companies)


	transaction	client_type
0	transac_olp	veggie
1	transac_ixh	veggie
2	transac_qkh	vegan
3	transac_qlz	omnivorous
4	transac_mal	vegan
5	transac_fjh	veggie
6	transac_rdn	veggie
7	transac_oaj	omnivorous
8	transac_taz	vegan
9	transac_tgs	veggie



Multivariate Analysis

- Quantitative variables with linear relation
 - Scatter plot
 - Linear regression
 - Relation / link


	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17



Multivariate Analysis

- Quantitative variables with linear relation
 - Bar plot raw data
 - More interesting when Time Series (not here)

	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17



Multivariate Analysis

- Quantitative variables with linear relation
 - Correlation matrix
 - Pearson correlation (linear)
 - Symetric matrix

	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17


```
# Pearson correlation for linear relations  
df_multi_quantitative.corr()
```

	rain	temperature
rain	1.000000	-0.826323
temperature	-0.826323	1.000000

Multivariate Analysis

- Quantitative variables with linear relation
 - Pairplot: Histograms and scatter plots
 - Very useful for +3 variables

	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17



Multivariate Analysis

- Quantitative variables with linear relation
 - Groupby
 - Aggregation functions: mean, max, std, etc.
 - Groups: type of clients, of investments, etc.
 - Less information, more lisibility

	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17

```
groups = df_multi_quantitative.groupby('rain').mean()  
groups
```


✓ 0.2s

temperature	
rain	
0	20.75
50	19.00
80	20.00
100	17.00
120	16.00
150	17.00

Multivariate Analysis

- Quantitative variables with linear relation
 - Bar plot
 - Here based on grouped data
 - Complementary to raw data bar plot

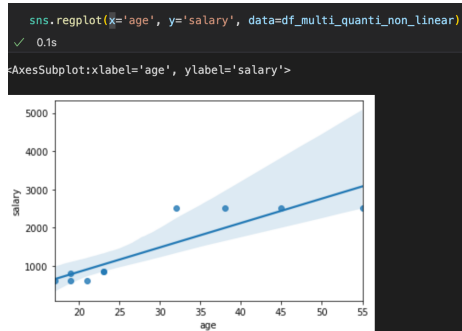
	rain	temperature
0	120	15
1	120	17
2	80	20
3	0	20
4	50	19
5	100	17
6	0	20
7	0	20
8	0	23
9	150	17



Multivariate Analysis

- Quantitative variables with non-linear relation
 - Scatter plot
 - Linear regression misleading

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850



Multivariate Analysis

- Quantitative variables with non-linear relation
 - Correlation matrix
 - Pearson correlation (linear) also misleading

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850

```
df_multi_quanti_non_linear.corr()
```

✓ 0.2s

	age	salary
age	1.000000	0.891305
salary	0.891305	1.000000

Multivariate Analysis

- Quantitative variables with non-linear relation
 - Correlation matrix
 - Spearman correlation (non-linear) instead
 - Rank based correlation

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850

```
from scipy import stats
stats.spearmanr(df_multi_quantitative_non_linear)
```

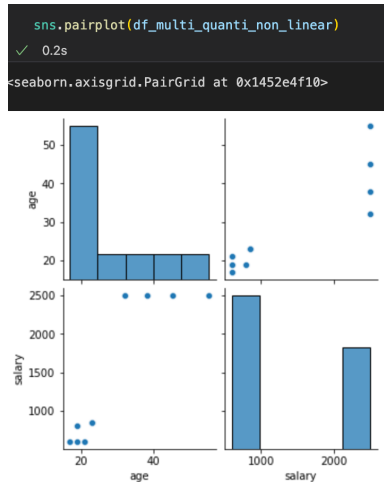
✓ 0.2s

SpearmanrResult(correlation=0.9209224503231543, pvalue=0.0001553715621233412)

Multivariate Analysis

- Quantitative variables with non-linear relation
 - Pairplot
 - Different regims well separated?

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850



Multivariate Analysis

- Quantitative variables with non-linear relation
 - Analyse regimes relations separately

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850



Multivariate Analysis

- Quantitative variables with non-linear relation
 - Analyse regimes correlations separately

	age	salary
0	17	600
1	19	600
2	21	600
3	45	2500
4	38	2500
5	55	2500
6	19	800
7	32	2500
8	23	850
9	23	850

```
sorted.iloc[6:].corr()
```

✓ 0.4s

	age	salary
age	1.000000	0.682242
salary	0.682242	1.000000

```
sorted.iloc[6:].corr()
```

✓ 0.3s

	age	salary
age	1.0	NaN
salary	NaN	NaN

Multivariate Analysis

- Qualitative variables
 - Counts
 - Proportions on pie chart

	city	sweat_size
0	paris	M
1	marseille	S
2	lyon	M
3	marseille	S
4	paris	M
5	marseille	S
6	bordeaux	XS
7	montpellier	M
8	paris	M
9	paris	M

```
# Counts are useful to obtain a  
# quantitative variable  
df_multi_quali.value_counts()
```

✓ 0.4s


city	sweat_size	
paris	M	4
marseille	S	3
bordeaux	XS	1
lyon	M	1
montpellier	M	1

dtype: int64

```
df_multi_quali.value_counts().plot.pie()
```

✓ 0.6s

<AxesSubplot:ylabel='None'>



Multivariate Analysis

- Qualitative variables
 - Visualize modalities separately

	city	sweat_size
0	paris	M
1	marseille	S
2	lyon	M
3	marseille	S
4	paris	M
5	marseille	S
6	bordeaux	XS
7	montpellier	M
8	paris	M
9	paris	M

```
df_multi_quali.groupby('city').count()
```

✓ 0.3s

sweat_size	
city	
bordeaux	1
lyon	1
marseille	3
montpellier	1
paris	4

```
df_multi_quali.groupby('sweat_size').count()
```

✓ 0.2s

city	
sweat_size	
M	6
S	3
XS	1

Multivariate Analysis

- Qualitative variables
 - Counts on group by

	city	sweat_size
0	paris	M
1	marseille	S
2	lyon	M
3	marseille	S
4	paris	M
5	marseille	S
6	bordeaux	XS
7	montpellier	M
8	paris	M
9	paris	M

```
df_multi_quali.groupby('city').count().plot.pie(y='sweat_size')
```

✓ 0.8s


<AxesSubplot:ylabel='sweat_size'>



```
df_multi_quali.groupby('sweat_size').count().plot.pie(y='city')
```

✓ 0.7s

<AxesSubplot:ylabel='city'>



Multivariate Analysis

- Qualitative variables
 - Counts on group by

	city	sweat_size
0	paris	M
1	marseille	S
2	lyon	M
3	marseille	S
4	paris	M
5	marseille	S
6	bordeaux	XS
7	montpellier	M
8	paris	M
9	paris	M

```
# test correlation: is there a link between the sweat size and the city?
from scipy.stats import chi2_contingency
crosstab = pd.crosstab(df_multi_quali['city'], df_multi_quali['sweat_size'])
crosstab
```

✓ 0.4s

sweat_size	M	S	XS
city			
bordeaux	0	0	1
lyon	1	0	0
marseille	0	3	0
montpellier	1	0	0
paris	4	0	0

```
results = chi2_contingency(crosstab)
print('pvalue', results[1])
```

✓ 0.2s


pvalue 0.010336050675925726

Multivariate Analysis

- Quantitative with qualitative variable
 - Boxplot
 - Quality of separation between modalities

	age_company_months	default
0	13	1
1	100	1
2	90	1
3	300	0
4	12	1
5	20	1
6	500	0
7	40	1
8	1	1
9	250	0

```
sns.boxplot(x='default', y='age_company_months',  
            data=df_multi_quantitative)  
✓ 0.6s  
<AxesSubplot:xlabel='default', ylabel='age_company_mon
```



Multivariate Analysis

- Quantitative with qualitative variable
 - Statistical test ANOVA
 - Correlation

	city	sweat_size
0	paris	M
1	marseille	S
2	lyon	M
3	marseille	S
4	paris	M
5	marseille	S
6	bordeaux	XS
7	montpellier	M
8	paris	M
9	paris	M






```
# Verify this correlation using ANOVA test
from statsmodels.formula.api import ols
from statsmodels.api import stats
model = ols('default ~ age_company_months', data=df_multi_quanti_quali).fit()
anova_table = stats.anova_lm(model)
anova_table
```

✓ 0.3s

	df	sum_sq	mean_sq	F	PR(>F)
age_company_months	1.0	0.715502	0.715502	4.134365	0.076463
Residual	8.0	1.384498	0.173062	NaN	NaN

Correlation


- Move repetitively in conjunction
- Methods
 - Pearson
 - Spearman (Rank)
- Spurious correlation (ice cream, Eiffel Tower)



Statistical Laws


- Normal Law / Gauss Curve

- Totally resumed by mean and variance
- Constant mean (0 if centered) and variance (1 if reduced)
- Uncorrelated individuals
- Symetric (Skewness=0)
- Precise bell shape (Kurtosis=3)



Statistical Laws

- Power Laws: multiplicative growth
- Examples:
 - Normal: human age, size, weight, grades
 - Power: lakes size, wealth



Statistical Tests

- Intention: prevent sampling error
- Hypothesis (Normal Law)
- Examples:
 - Normality test
 - ANOVA
 - Pearson's r
 - Chi square

